

# LOGISTIC REGRESSION IN PYTHON

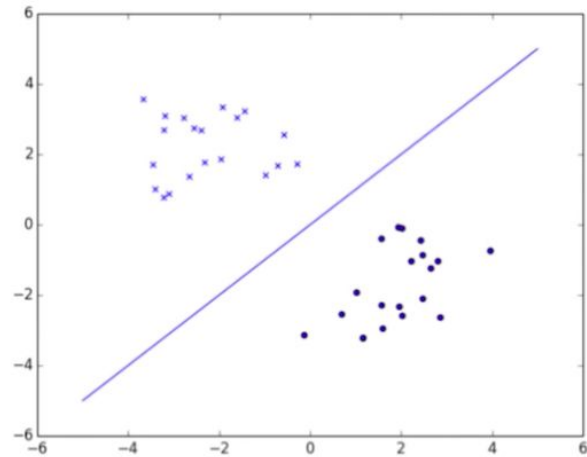
## 2-dimensional classification

$$y = mx + b$$

or

$$0 = ax + by + c$$

$$a = 1, b = -1, c = 0$$



## Machine learning lingo

We call  $(x, y) \rightarrow (x_1, x_2) = \mathbf{x}$

We rename the constants to  $w_i$

We call the bias term / intercept  $w_0$

$$h(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2$$

We say  $h()$  is a **linear combination** of the components of  $\mathbf{x}$

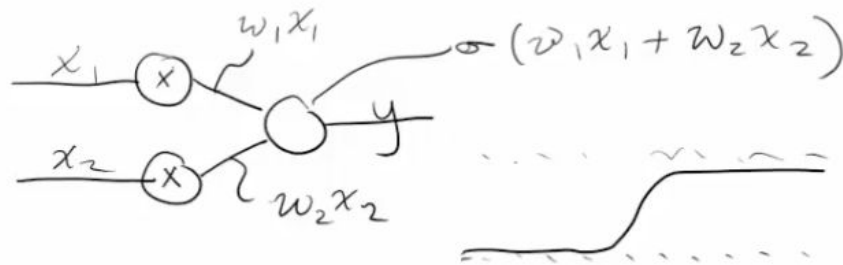
$$\text{In vector form: } h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$

In 3-dimensions: line  $\rightarrow$  plane, In  $> 3$  dimensions: hyperplane

# Neuron Analogies

1. Many inputs  $\rightarrow$  One output
2. Spike or no spike  $\rightarrow$  0/1 output
3. Synapse strengths  $\rightarrow$  linear weights

mber 6, 2015 11:24 PM



$$\tanh(x) \in (-1, 1)$$

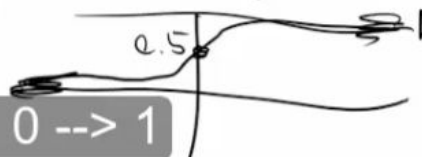
$y \rightarrow 0$

$$\sigma(z) = \frac{1}{1 + e^{-z}} \in (0, 1)$$

$y \rightarrow 0.5$

$$\sigma(w^T x)$$

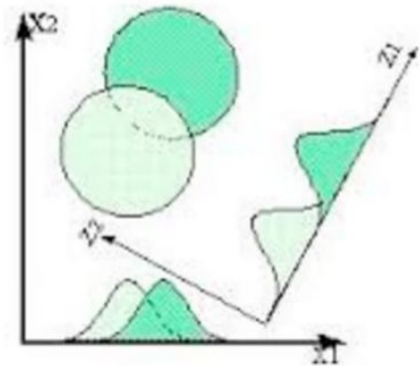
Before we had  $w^*x > 0 \rightarrow 1$



## Problem Description

- 2 Gaussian Clouds
- Both have the same covariance
- Multivariate Gaussian PDF:

$$p(x) = \frac{1}{\sqrt{(2\pi)^D |\Sigma|}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)}$$



## Bayes' Rule

$$p(Y | X) = p(X | Y)p(Y) / p(X)$$

i.e.:

$$p(Y=1 | X) = p(X | Y=1)p(Y=1) / p(X)$$

$$p(Y=0 | X) = p(X | Y=0)p(Y=0) / p(X)$$

- $p(X | Y)$  is the Gaussian - we calculate it over all the data that belongs to class  $Y$
- $p(Y)$  is just the frequency estimate of  $Y \rightarrow$   
e.g.  $p(Y=1) = \text{\#times class 1 appeared} / \text{\# total}$

## Put it into the logistic regression framework

Manipulate Bayes' Rule

$$p(y = 1 | x) = \frac{p(x | y=1)p(y=1)}{p(x)} = \frac{p(x | y=1)p(y=1)}{p(x | y=1)p(y=1) + p(x | y=0)p(y=0)}$$

Divide top and bottom by  $p(x | y=1)p(y=1)$

$$p(y = 1 | x) = \frac{1}{1 + \frac{p(x | y=0)p(y=0)}{p(x | y=1)p(y=1)}}$$

Looks a lot like Logistic Regression!

## Put it into the logistic regression framework

Manipulate Bayes' Rule

$$p(y = 1 | x) = \frac{p(x | y=1)p(y=1)}{p(x)} = \frac{p(x | y=1)p(y=1)}{p(x | y=1)p(y=1) + p(x | y=0)p(y=0)}$$

Divide top and bottom by  $p(x | y=1)p(y=1)$

$$p(y = 1 | x) = \frac{1}{1 + \frac{p(x | y=0)p(y=0)}{p(x | y=1)p(y=1)}}$$

Looks a lot like Logistic Regression!

## Logistic Regression

$$p(y = 1 \mid x) = \frac{1}{1 + \frac{p(x \mid y=0)p(y=0)}{p(x \mid y=1)p(y=1)}} = \frac{1}{1 + \exp(-w^T x)}$$

$$-w^T x = \ln \left( \frac{p(x \mid y=0)p(y=0)}{p(x \mid y=1)p(y=1)} \right)$$

Let's just manipulate the right side and remember that it equals to the left

## Logistic Regression

Let:  $p(y = 1) = 1 - \alpha$ ,  $p(y = 0) = \alpha$

$$\ln \left( \frac{p(x \mid y=0)p(y=0)}{p(x \mid y=1)p(y=1)} \right) = \ln p(x \mid y = 0) + \ln \alpha - \ln p(x \mid y = 1) - \ln(1 - \alpha)$$

$$= \ln \frac{1}{\sqrt{(2\pi)^D |\Sigma|}} e^{-\frac{1}{2}(x-\mu_0)^T \Sigma^{-1}(x-\mu_0)} - \ln \frac{1}{\sqrt{(2\pi)^D |\Sigma|}} e^{-\frac{1}{2}(x-\mu_1)^T \Sigma^{-1}(x-\mu_1)} + \ln \frac{\alpha}{1-\alpha}$$

## Logistic Regression

$$= \ln \frac{1}{\sqrt{(2\pi)^D |\Sigma|}} e^{-\frac{1}{2}(x-\mu_0)^T \Sigma^{-1} (x-\mu_0)} - \ln \frac{1}{\sqrt{(2\pi)^D |\Sigma|}} e^{-\frac{1}{2}(x-\mu_1)^T \Sigma^{-1} (x-\mu_1)} + \ln \frac{\alpha}{1-\alpha}$$

1 / sqrt(...) cancels out

$$\begin{aligned} = & -\frac{1}{2}(x^T \Sigma^{-1} x - x^T \Sigma^{-1} \mu_0 - \mu_0^T \Sigma^{-1} x + \mu_0^T \Sigma^{-1} \mu_0) \\ & + \frac{1}{2}(x^T \Sigma^{-1} x - x^T \Sigma^{-1} \mu_1 - \mu_1^T \Sigma^{-1} x + \mu_1^T \Sigma^{-1} \mu_1) \\ & + \ln \frac{\alpha}{1-\alpha} \end{aligned}$$

## Logistic Regression

$$\begin{aligned} = & -\frac{1}{2}(x^T \Sigma^{-1} x - x^T \Sigma^{-1} \mu_0 - \mu_0^T \Sigma^{-1} x + \mu_0^T \Sigma^{-1} \mu_0) \\ & + \frac{1}{2}(x^T \Sigma^{-1} x - x^T \Sigma^{-1} \mu_1 - \mu_1^T \Sigma^{-1} x + \mu_1^T \Sigma^{-1} \mu_1) \\ & + \ln \frac{\alpha}{1-\alpha} \end{aligned}$$

- Quadratic term cancels out
- Is  $x^T \Sigma^{-1} \mu = \mu^T \Sigma^{-1} x$ ? Yes! Try it on paper to prove to yourself. Remember the covariance, and hence its inverse, is symmetric

## Logistic Regression

$$= \mu_0^T \Sigma^{-1} x - \mu_1^T \Sigma^{-1} x - \frac{1}{2} \mu_0^T \Sigma^{-1} \mu_0 + \frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 + \ln \frac{\alpha}{1-\alpha}$$

$$= (\mu_0^T - \mu_1^T) \Sigma^{-1} x - \frac{1}{2} \mu_0^T \Sigma^{-1} \mu_0 + \frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 + \ln \frac{\alpha}{1-\alpha}$$

$$= -(w^T x + b)$$

Finally!

$$w^T = (\mu_1^T - \mu_0^T) \Sigma^{-1}$$

$$b = \frac{1}{2} \mu_0^T \Sigma^{-1} \mu_0 - \frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 - \ln \frac{\alpha}{1-\alpha}$$

Our data

$$\mu_0 = (-2, -2)^T, \mu_1 = (+2, +2)^T, \sigma^2 = 1$$

$w = (4, 4)$ ,  $b = 0$  (assume equal number of samples of each class)

- Exercise: Solve it yourself on paper!
- Exercise: Make sure you can do the derivation yourself without watching the lectures.
- Also called: Linear Discriminant Analysis (LDA)
- If covariance is a diagonal matrix: Naive Bayes (why?)
- If we have different covariances: Quadratic Discriminant Analysis (QDA)
  - Exercise: Write your own code and compare the performance vs. LDA
- This solution is optimal provided assumptions about distributions are true



## Input Size

- $N$  = number of samples we've collected
  - i.e. I measure the height of 10 students to find the average height.  $N=10$
- $D$  = number of dimensions or features per sample
  - i.e. I measure height, weight, and girth, to try and predict body fat %.  $D=3$  (height, weight, girth)
- A matrix of data called  $X$  would be an  $N \times D$  matrix
  - Each row is a sample
  - Each column is the value of one feature in each sample

## Data and Targets

- Sometimes we say we train a model on inputs  $X$  and targets  $Y$
- If we consider all  $Y$ s at the same time, then it is an  $N \times 1$  matrix (for binary classification, an  $N \times 1$  matrix of 0s and 1s)
- Sometimes we use  $T$  for the target instead
  - i.e. The cross-entropy error function:  $-[t \log(y) + (1 - t) \log(1 - y)]$
- Now  $Y$  is used for something else: output of logistic regression.
- Before the output was more precise,  $P(Y=1 | X)$
- But that takes longer to write...

## Cost function

- Cost function = error function = objective function
- We want to minimize cost and error - easy to see semantically
- Objective function can mean something to maximize OR minimize
- Trivial sign flipping - minimizing  $x^2$  = maximizing  $-x^2$
- We will first maximize likelihood -  $P(\text{data} | \text{model})$
- Same as maximizing log-likelihood  $L = \log P(\text{data} | \text{model})$ 
  - i.e. If  $A > B$ , then  $\log(A) > \log(B)$
- Sometimes, we use the letter  $J$ 
  - May be in the form  $J = L \rightarrow$  Maximize it
  - Or  $J = -L \rightarrow$  Minimize it



# Logistic Regression Error

Linear Regression - Squared Error

$$J = \sum_n (t_n - y_n)^2$$

Assumes Gaussian-distributed error, because  $\log(\text{Gaussian}) = \text{squared function}$

Logistic Regression Error can't be Gaussian distributed, because:

- Target is only 0/1
- Output is only a number between 0-1

We want: 0 if correct, > 0 if not correct, more wrong == bigger cost

# Cross-Entropy Error

$$J = -\{ t \log(y) + (1-t) \log(1-y) \} \quad (t = \text{target}, y = \text{output of logistic})$$

If  $t = 1$ , only first term matters, if  $t = 0$ , only second term matters

$\log(y) \rightarrow$  number between 0 and  $-\infty$

Ex:

- $t = 1, y = 1 \rightarrow 0$
- $t = 0, y = 0 \rightarrow 0$
- $t = 1, y = 0.9 \rightarrow 0.11$
- $t = 1, y = 0.5 \rightarrow 0.69$
- $t = 1, y = 0.1 \rightarrow 2.3$

## Multiple Training Examples

$$J = - \sum_{n=1}^N t_n \log(y_n) + (1 - t_n) \log(1 - y_n)$$

Saturday, November 7, 2015

12:18 AM

MAXIMUM LIKELIHOOD

can  $P(H) = p$   $P(T) = 1 - p$   
 $N = 10$  7H 3T

$$L = p^7 (1-p)^3$$

$$\begin{aligned} \ell &= \log [p^7 (1-p)^3] \\ &= \log p^7 + \log (1-p)^3 \\ &= 7 \log p + 3 \log (1-p) \end{aligned}$$

$$L = \log [p^7 (1-p)^3]$$

$$\begin{aligned} l &= \log [p^7 (1-p)^3] \\ &= \log p^7 + \log (1-p)^3 \\ &= 7 \log p + 3 \log (1-p) \end{aligned}$$

$$\frac{\partial l}{\partial p} = \frac{7}{p} + \frac{3}{1-p} (-1) = 0$$

$$\frac{7}{p} = \frac{3}{1-p}$$

$$\frac{1-p}{p} = \frac{3}{7}$$

$$\begin{aligned}
 \ell &= \log [p^7 (1-p)^3] \\
 &= \log p^7 + \log (1-p)^3 \\
 &= 7 \log p + 3 \log (1-p)
 \end{aligned}$$

$$\frac{\partial \ell}{\partial p} = \frac{7}{p} + \frac{3}{1-p} (-1) = 0$$

$$\frac{7}{p} = \frac{3}{1-p}$$

$$\frac{1-p}{p} = \frac{3}{7}$$

$$\frac{1}{p} - 1 = \frac{3}{7}$$

$$\frac{1}{p} = \frac{10}{7}$$

$$p = \frac{7}{10} = p^{(H)}$$

$$p(y=1|x) = \sigma(w^T x) = y$$

$$L = \prod_{n=1}^N y_n^{t_n} (1-y_n)^{1-t_n}$$

$$\ell = \sum_n t_n \log y_n + (1-t_n) \log (1-y_n)$$

# How to optimize the weights

- IF we assume data is Gaussian-distributed with equal covariance, we can use the Bayes method
- But we want something that will work in general
- Linear regression:
- Take derivative, set to 0, solve for weights
- Can't do this with logistic regression / cross-entropy (but you're encouraged to try)

## Gradient Descent

Idea: take small steps in direction of derivative

Step size == learning rate (1 for this example)

Ex:

$$w = -2$$

$$w = -2 - 1 * (-1) = -1$$

Now we're closer to the optimal point! ( $w=0$ )

Note: slope is 0 at the bottom, so no more changes will occur



## Bias Term

$$\frac{\partial J}{\partial w_0} = \sum_{n=1}^N (y_n - t_n) x_{n0} = \sum_{n=1}^N (y_n - t_n)$$

# Gradient Descent for Logistic Regression

$$J = - \sum_{n=1}^N t_n \log(y_n) + (1 - t_n) \log(1 - y_n)$$

Split into 3 derivatives:

$$\frac{\partial J}{\partial w_i} = \sum_{n=1}^N \frac{\partial J}{\partial y_n} \frac{\partial y_n}{\partial a_n} \frac{\partial a_n}{\partial w_i}$$
$$a_n = w^T x_n$$

## Derivatives

$$J = - \sum_{n=1}^N t_n \log(y_n) + (1 - t_n) \log(1 - y_n)$$

$$\frac{\partial J}{\partial y_n} = - t_n \frac{1}{y_n} + (1 - t_n) \frac{1}{1 - y_n} (-1)$$



## Derivatives

$$y_n = \sigma(a_n) = \frac{1}{1+e^{-a_n}}$$

$$\frac{\partial y_n}{\partial a_n} = \frac{-1}{(1+e^{-a_n})^2} (e^{-a_n})(-1)$$

$$\frac{\partial y_n}{\partial a_n} = \frac{e^{-a_n}}{(1+e^{-a_n})^2} = \frac{1}{1+e^{-a_n}} \frac{e^{-a_n}}{1+e^{-a_n}} = y_n(1 - y_n)$$

## Derivatives

$$a_n = w^T x_n$$

$$a_n = w_0 x_{n0} + w_1 x_{n1} + w_2 x_{n2} + \dots$$

$$\frac{\partial a_n}{\partial w_i} = x_{ni}$$

## Putting them all together

$$\frac{\partial J}{\partial w_i} = - \sum_{n=1}^N \frac{t_n}{y_n} y_n (1 - y_n) x_{ni} - \frac{1-t_n}{1-y_n} y_n (1 - y_n) x_{ni}$$

$$\frac{\partial J}{\partial w_i} = - \sum_{n=1}^N t_n (1 - y_n) x_{ni} - (1 - t_n) y_n x_{ni}$$

$$\frac{\partial J}{\partial w_i} = - \sum_{n=1}^N [t_n - t_n y_n - y_n + t_n y_n] x_{ni}$$

$$\frac{\partial J}{\partial w_i} = \sum_{n=1}^N (y_n - t_n) x_{ni}$$

# Vectorize

$$\frac{\partial J}{\partial w} = \sum_{n=1}^N (y_n - t_n) x_n$$

## Vectorize More

Dot product is a sum over some index

X is N x D

Y, T are N x 1

$$a^T b = \sum_{n=1}^N a_n b_n$$

Multiply  $X^T(Y - T)$ :

Shape is (D x N)(N x 1)  $\rightarrow$  (D x 1)

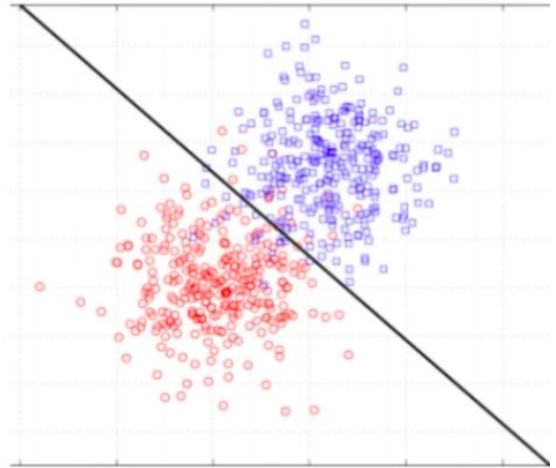
Which is the correct shape for w

N gets summed over

$$\frac{\partial J}{\partial w} = X^T (Y - T)$$

## Generalization and Overfitting

- Recall our “Gaussian cloud” problem
- 2 clouds
- One centered at (2,2)
- Other centered at (-2,-2)
- Exact Bayesian solution was:
- $w = [0, 4, 4]$
- Represent as  $y = mx + b$
- High school math!
- $0 + 4x + 4y = 0$
- $y = -x$



$$y = -x$$

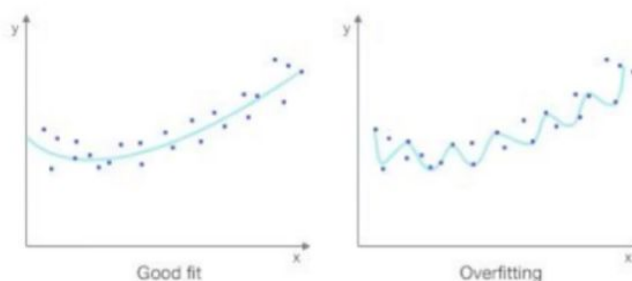
- Slope: -1, y-intercept: 0
- Why is the Bayesian solution (4, 4)?
- Why not (1, 1)? Why not (10, 10)?
- These all would represent the same line
- First hint at why we might need regularization

## Objective function

- $J = t \log(y) + (1-t) \log(1-y)$
- This  $y$  is the output of logistic regression
- Take a test point  $(x_1, x_2) = (1, 1)$  and existing weights  $(0, 4, 4)$
- Should be classified as 1
- $\sigma(0 + 4 + 4) = \sigma(8) = 0.99966$
- What would be better? Exactly 1!
- With  $y = \sigma(8)$ ,  $J = -0.00033540637289566265$
- What if my weights are  $(0, 1, 1)$ ?  $\rightarrow J = -0.12692801104297263$
- Not as good
- What if my weights are  $(0, 10, 10)$ ?  $\rightarrow J = -2.0611536942919273e-09$
- The “best” weights are thus  $(0, \text{infinity}, \text{infinity})$
- In the computer, that’s an error

## Regularization

- People usually explain regularization in terms of overfitting (image on the right)
- But this is regression so it’s not exactly applicable
- If your data fills up the space of all possible inputs, you shouldn’t overfit even if your model is very complex
- That’s why we want to have lots of data
- Your model overfits when it has to “guess” what the output should be in a space it’s never seen
- But if the train set entirely covers all future test sets, then it’s fine

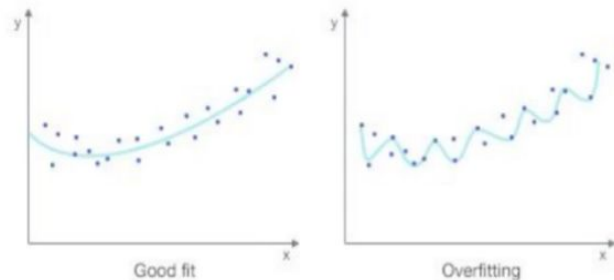


# Regularization

- This scenario is different - we could have a perfectly split up dataset that covers the entire possible input space
- Logistic regression would still try to go to  $w = (0, \text{infinity}, \text{infinity})$
- Solution: regularization
- Penalizes very large weights
- Existing cost function:  $J = -[t \log(y) + (1 - t) \log(1 - y)]$
- Add a penalty for big weights:
- $J_{\text{reg}} = J + (\lambda/2) \|w\|^2 = J + (\lambda/2) w^T w$
- Now we won't try to go to  $(0, 10, 10)$
- $\lambda$  = smoothing parameter (usually  $\sim 0.1, 1$ ) but depends on your data
- No universal formula to choose (maybe, but too advanced for this class)

## Regularization

- People usually explain regularization in terms of overfitting (image on the right)
- But this is regression so it's not exactly applicable
- If your data fills up the space of all possible inputs, you shouldn't overfit even if your model is very complex
- That's why we want to have lots of data
- Your model overfits when it has to "guess" what the output should be in a space it's never seen
- But if the train set entirely covers all future test sets, then it's fine



# Regularization

- This scenario is different - we could have a perfectly split up dataset that covers the entire possible input space
- Logistic regression would still try to go to  $w = (0, \text{infinity}, \text{infinity})$
- Solution: regularization
- Penalizes very large weights
- Existing cost function:  $J = -[t \log(y) + (1 - t) \log(1 - y)]$
- Add a penalty for big weights:
- $J_{\text{reg}} = J + (\lambda/2) \|w\|^2 = J + (\lambda/2) w^T w$
- Now we won't try to go to  $(0, 10, 10)$
- $\lambda$  = smoothing parameter (usually  $\sim 0.1, 1$ ) but depends on your data
- No universal formula to choose (maybe, but too advanced for this class)

## Solving for w

- Still do gradient descent
- Now we want  $dJ_{\text{reg}} / dw$  instead of  $dJ / dw$
- Addition doesn't affect other terms in derivative
- $\text{reg\_cost} = (\lambda/2) (w_0^2 + w_1^2 + w_2^2 + \dots)$
- $d(\text{reg\_cost})/dw_i = \lambda w_i$
- In vector form:  $d(\text{reg\_cost})/dw = \lambda w$
- $dJ_{\text{reg}} / dw = X^T(Y - T) + \lambda w$



## Probabilistic Perspective

- Cross-entropy maximizes the likelihood, since  $J = -\log(\text{likelihood})$
- Let's make  $J$  something we want to maximize:
- $J = +[t \log y + (1 - t) \log(1 - y)] - (\lambda/2)\|w\|^2$
- $\exp(J) = y^t(1 - y)^{(1 - t)} \exp(-\lambda\|w\|^2 / 2)$
- Exponentiate the first part: Bernoulli (the likelihood)
- Exponentiate -ve squared term: Gaussian (the prior)
- $t \sim \text{Bernoulli}(y)$
- $w \sim N(0, 1/\lambda)$

## Probabilistic Perspective

- We have one probability distribution x another probability distribution
- Your first baby steps toward a Bayesian perspective of ML
- $\text{posterior} \propto \text{likelihood} \times \text{prior}$
- $t \sim \text{Bernoulli}(y) \rightarrow (\text{likelihood})$
- $w \sim N(0, 1/\lambda) \rightarrow (\text{prior})$
- Our “prior belief” about  $w$  is that it's Gaussian distributed with variance  $1/\lambda$

## Bayes rule

- $P(w|X, Y) \propto P(X, Y|w) P(w)$
  - Just a consequence of:
  - $P(B|A) = P(A|B)P(B) / P(A) = P(A,B) / P(A)$
  - Bottom part is just  $\sum_B P(A,B) = \sum_B P(A|B)P(B)$
- 
- Without regularization, we maximize the likelihood
  - With regularization, we maximize the posterior
  - This is called “maximum a posteriori” or MAP estimation

## L1 Regularization

- Previously: you saw that even adding a column of completely random noise can improve  $R^2$
- In general, we want  $D \ll N$  (# features  $\ll$  # samples)

Skinny  
( $D \ll N$ )  
(ideal)

Fat ( $D > N$ ) (not ideal)

# L1 Regularization

- L2 regularization used L2 norm for penalty term
- As you might guess, L1 regularization uses L1 norm for penalty term

$$J_{RIDGE} = - \sum_{n=1}^N (t_n \log y_n + (1 - t_n) \log(1 - y_n)) + \lambda \|w\|_2^2$$

$$J_{LASSO} = - \sum_{n=1}^N (t_n \log y_n + (1 - t_n) \log(1 - y_n)) + \lambda \|w\|_1$$

# L1 Regularization

- This also puts a prior on  $w$ , so it's also a MAP estimation of  $w$
- We can determine the distribution by taking the negative exponential of the penalty
- Laplace distribution

$$p(w) = \frac{\lambda}{2} \exp(-\lambda |w|)$$

# L1 Regularization

- This also puts a prior on  $w$ , so it's also a MAP estimation of  $w$
- We can determine the distribution by taking the negative exponential of the penalty
- Laplace distribution

$$p(w) = \frac{\lambda}{2} \exp(-\lambda|w|)$$

# L1 Regularization

$$J = - \sum_{n=1}^N t_n \log(y_n) + (1 - t_n) \log(1 - y_n) + \lambda|w|$$

$$\partial J / \partial w = X^T (Y - T) + \lambda \text{sign}(w)$$

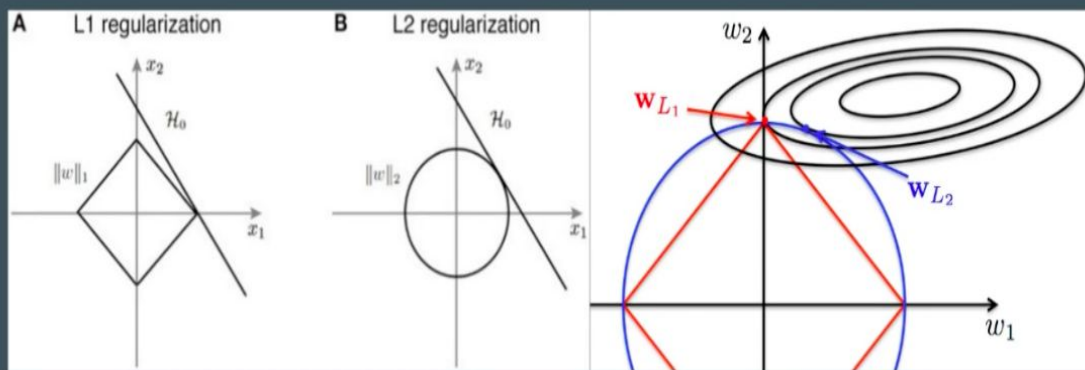
- $\text{sign}(x) = 1$  if  $x > 0$ ,  $-1$  if  $x < 0$ ,  $0$  if  $x = 0$
- Use the same gradient descent algorithm we've already been using

# L1 vs L2 Regularization

- What's the difference?
  - We'll look at the result first, then at how they emerge from the math
  - L1: encourages a sparse solution (few  $w$ 's non-zero, many equal to 0)
  - L2: encourages small weights (all  $w$ 's close to 0, but not exactly 0)
- 
- Both help you prevent overfitting, by not fitting to noise
  - L1 accomplishes this by choosing the most important features
  - L2 accomplishes this by making the assertion that none of the weights are extremely large

## L1 vs L2 Regularization

- I find this visualization awful, yet it's very commonly taught
- Contour plot of negative log-priors
- Ugh!



# L1 vs L2 Regularization

Think of a 1-dimensional weight:

L2 penalty is quadratic, L1 is absolute function

Key is the derivative (since the solution uses gradient descent)

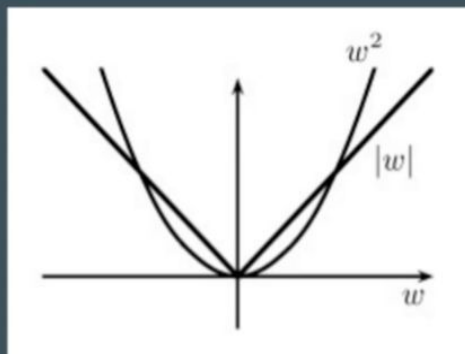
Quadratic: as  $w \rightarrow 0$ , derivative  $\rightarrow 0$

If  $w$  is already small, further gradient descent won't change it much

Absolute: derivative is always  $\pm 1$  (0 at  $w=0$ )

Doesn't matter where  $w$  is, it will fall at a constant rate

When it reaches 0, it stays there forever!



## ElasticNet

- It's possible to include both L1 and L2 simultaneously (fancy name is ElasticNet)

$$J_{RIDGE} = J + \lambda_2 |w|^2$$

$$J_{LASSO} = J + \lambda_1 |w|$$

$$J_{ELASTICNET} = J + \lambda_1 |w| + \lambda_2 |w|^2$$

## Outline of our sentiment analyzer

- We'll just look at the electronics category, but you can try the same code on others
- We could use 5 star targets to do regression, but let's just do classification since they are already marked "positive" and "negative"
- XML parser (BeautifulSoup)
- Only look at key "review\_text"
- We'll need 2 passes, one to determine vocabulary size and which index corresponds to which word, and one to create data vectors
- After that, we can just use any SKLearn classifier as we did previously
- But we'll use logistic regression so we can interpret the weights

**Skipping implementation details..**

**$X$  = one-hot encoded bag of words**

**$Y$  = 1/0 (positive/negative)**



## 2-class problem vs. 7-class problem

- When we switch to softmax (which we will in Deep Learning part 1), will the problem get easier or harder?
- 2 class:
  - Guess at random - expect 50% error
- 7 class:
  - Guess at random - expect  $6/7 = 86\%$  error
- K class:  $1/K$  chance of being correct

Kaggle top score: ~70% correct

## Normalize the data

- Images have pixel intensities 0... 255 (8 bit integers have  $2^8 = 256$  different possible values)
- We want to normalize these to be from 0... 1
- (Another way to normalize is  $z = (x - \text{mean}) / \text{stddev}$ )
- Reason: sigmoid / tanh are most active in the -1... +1 range

## Sensitivity and Specificity

- Used in medical field
- Sensitivity:

True positive rate = TPR =  $TP / (TP + FN)$

- Specificity:

True negative rate = TNR =  $TN / (TN + FP)$

# Precision and Recall

- Used in information retrieval
- Precision

$$TP / (TP + FP)$$

- Recall

$$TP / (TP + FN)$$

- Note: recall = sensitivity

## F1-score

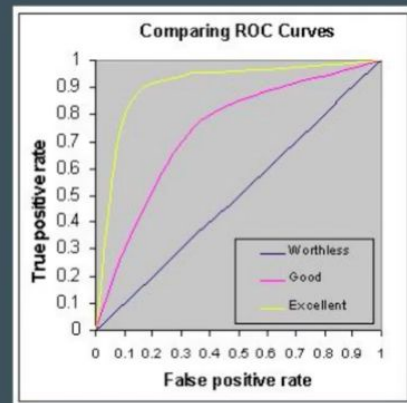
- Combines precision and recall into a balanced measure

$$F1 = 2 * (precision * recall) / (precision + recall)$$

- Harmonic mean of precision and recall

## ROC and AUC

- In logistic regression we use 0.5 as a threshold, makes sense because  $P(Y=1|X) > 0.5 \rightarrow \text{guess } 1$ ,  $P(Y=1|X) < 0.5 \rightarrow \text{guess } 0$
- But we can use any threshold
- Different thresholds  $\rightarrow$  different TPR and FPR
- Receiver operating characteristic (ROC curve is a plot of these for every value of threshold between 0 and 1)
- Area under curve = AUC
  - 1 = perfect classifier, 0.5 = random guessing

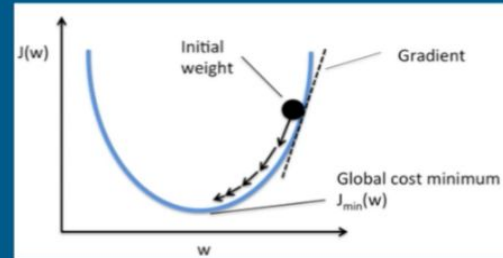


# Unsupervised Machine Learning

Hidden Markov Models in Python  
LazyProgrammer

# Gradient Descent Tutorial

- Optimization method
- Used extensively in deep learning, useful in a wide variety of situations
- Idea:
- You have a function you want to minimize,  $J(w)$  = cost or error
- Can maximize things too, just switch signs



## Example

$$J = w^2$$

(we know min is at  $w = 0$ , but let's pretend we don't)

$dJ/dw = 2w$ , set initial  $w = 20$ , learning rate = 0.1

Iteration 1:  $w \leftarrow 20 - 0.1 \cdot 40 = 16$

Iteration 2:  $w \leftarrow 16 - 0.1 \cdot 2 \cdot 16 = 12.8$

Iteration 3:  $w \leftarrow 12.8 - 0.1 \cdot 2 \cdot 12.8 = 10.24$

## Why is it so important?

- As we progress in deep learning / machine learning, functions will get more complicated
- Regular neural networks with softmax → might take you few hours or days to get derivatives the first time
  - Hopefully you did your homework!
- Convolutional and Recurrent Networks → possible, but don't want to
- Theano / TensorFlow can calculate gradients for us
- Understand how it works, so you can use it on anything

## Exercise

Try to optimize:  $J(w_1, w_2) = w_1^2 + w_2^4$

Maximizing the likelihood is the same as minimizing the cross-entropy function.